

CS6.302 - Software System Development

Lab Activity: The Performance Cost of Unbalanced Trees

Topic: BST vs. AVL Tree Performance Analysis

Scenario: The Network Anomaly Detector

You are a software engineer at a cybersecurity firm. Your team is building a high-speed anomaly detector that logs millions of network events. Each event has a unique, integer-based “event ID”. You need to store these IDs in a way that allows for two critical operations:

1. **Fast Insertion:** New events arrive constantly and must be logged.
2. **Ultra-Fast Lookup:** A separate process must constantly query the system to see if a specific event ID has already been logged.

A junior developer implemented a standard Binary Search Tree (BST). In testing with random data, performance was great ($O(\log n)$). However, when the system went live, it monitoring a specific subnet that was compromised. This subnet sent a massive flood of packets with sequential, increasing event IDs (e.g., 1001, 1002, 1003...).

The system’s performance immediately collapsed. Your job is to demonstrate **why** this happened and implement a robust solution (an AVL Tree) that can handle this “worst-case” data.

Your Task: An Automated Performance Analysis

You will implement both a standard **BST** and a self-balancing **AVL Tree**. You will feed both trees a “killer” dataset of 50,000 sorted integers read from a file. Finally, you will time their search performance and generate a report.

This is an auto-graded assignment. Your final script must produce a file named `analysis.txt` in the exact format specified.

Step 1: Define the ‘Node’ Class

Create a file named `avl_vs_bst.py`. Use this provided Node class.

```
1 from dataclasses import dataclass
2 from typing import Optional
3
4 @dataclass
5 class Node:
6     key: int
7     height: int = 1
8     left: Optional['Node'] = None
9     right: Optional['Node'] = None
```

Step 2: Implement the 'BST' Class

In the same file, implement a standard Binary Search Tree. You must fill in the logic for all the methods below. All methods must be recursive (where appropriate).

```
1 class BST:
2     def __init__(self):
3         self.root: Optional[Node] = None
4
5     def insert(self, key: int):
6         self.root = self._insert(self.root, key)
7
8     def _insert(self, root: Optional[Node], key: int) -> Node:
9         # TODO: Implement recursive BST insertion
10
11     def search(self, key: int) -> bool:
12         return self._search(self.root, key)
13
14     def _search(self, root: Optional[Node], key: int) -> bool:
15         # TODO: Implement recursive BST search
16
17     def get_height(self) -> int:
18         return self._get_height(self.root)
19
20     def _get_height(self, root: Optional[Node]) -> int:
21         # TODO: Implement recursive height calculation
22
23     def is_balanced(self) -> bool:
24         return self._is_balanced(self.root)
25
26     def _is_balanced(self, root: Optional[Node]) -> bool:
27         # TODO: Implement recursive balance check.
```

Step 3: Implement the 'AVLTree' Class

Now, implement the AVL tree. It will be very similar to the BST, but with crucial differences in the insert method and new helper functions.

```
1 class AVLTree:
2     def __init__(self):
3         self.root: Optional[Node] = None
4
5     # TODO: Implement _get_node_height (you can copy from BST, but it's
6     # simpler: return node.height if node else 0)
7     def _get_node_height(self, node: Optional[Node]) -> int:
8
9     # TODO: Implement get_balance
10    def _get_balance(self, node: Optional[Node]) -> int:
11
12    # TODO: Implement right_rotate
13    def _right_rotate(self, y: Node) -> Node:
14
15    # TODO: Implement left_rotate
16    def _left_rotate(self, x: Node) -> Node:
17
18    def insert(self, key: int):
19        self.root = self._insert(self.root, key)
20
21    def _insert(self, root: Optional[Node], key: int) -> Node:
22        # TODO: Implement recursive AVL insertion
23
24    def search(self, key: int) -> bool:
25        return self._search(self.root, key)
```

```

26
27 # TODO: Implement _search (can copy from BST)
28 def _search(self, root: Optional[Node], key: int) -> bool:
29     # ... (Same as BST's search)
30
31 def get_height(self) -> int:
32     return self._get_node_height(self.root)
33
34 def is_balanced(self) -> bool:
35     return self._is_balanced(self.root)
36
37 # TODO: Implement _is_balanced
38 def _is_balanced(self, root: Optional[Node]) -> bool:

```

Step 4: The Main Analysis Block

This is the most important part for grading. You **must** implement the main execution block to perform the analysis and write the `analysis.txt` file. You will be given `search_target.txt` and `sorted_sensor_data.txt`.

```

1 import time
2
3 if __name__ == "__main__":
4
5     data_filename = "sorted_sensor_data.txt"
6     target_filename = "search_target.txt"
7     output_filename = "analysis.txt"
8
9     # --- Load Data ---
10    print(f"Loading data from {data_filename}...")
11    keys_to_insert = []
12    try:
13        with open(data_filename, 'r') as f:
14            for line in f:
15                keys_to_insert.append(int(line.strip()))
16
17        with open(target_filename, 'r') as f:
18            search_key = int(f.read().strip())
19
20        print(f"Data loaded. {len(keys_to_insert)} keys.")
21        print(f"Search target: {search_key}")
22
23    except FileNotFoundError:
24        print(f"Error: Data files not found.")
25        exit()
26
27    import sys
28    # We must increase Python's recursion limit to allow the BST
29    # to build its extremely deep, skewed tree (height = 50,000).
30    # The default limit is ~1000.
31    new_limit = len(keys_to_insert) + 10
32    print(f"Setting recursion limit to {new_limit}...")
33    sys.setrecursionlimit(new_limit)
34
35    bst_tree = BST()
36    avl_tree = AVLTree()
37
38    # --- Time BST Insertion ---
39    print("Inserting into BST...")
40    bst_start_time = time.perf_counter()
41    for key in keys_to_insert:
42        bst_tree.insert(key)

```

```

43 bst_end_time = time.perf_counter()
44 bst_insert_time = bst_end_time - bst_start_time
45 print(f"BST insertion time: {bst_insert_time:.6f}s")
46
47 # --- Time AVL Insertion ---
48 print("Inserting into AVL...")
49 avl_start_time = time.perf_counter()
50 for key in keys_to_insert:
51     avl_tree.insert(key)
52 avl_end_time = time.perf_counter()
53 avl_insert_time = avl_end_time - avl_start_time
54 print(f"AVL insertion time: {avl_insert_time:.6f}s")
55
56 # --- 1. Get Tree Heights ---
57 bst_height = bst_tree.get_height()
58 avl_height = avl_tree.get_height()
59
60 # --- 2. Check Balance ---
61 bst_balanced = bst_tree.is_balanced()
62 avl_balanced = avl_tree.is_balanced()
63
64 # --- 3. Time Search ---
65 print(f"Searching for {search_key}...")
66
67 # Time BST Search
68 bst_search_start = time.perf_counter()
69 bst_tree.search(search_key)
70 bst_search_end = time.perf_counter()
71 bst_search_time = bst_search_end - bst_search_start
72
73 # Time AVL Search
74 avl_search_start = time.perf_counter()
75 avl_tree.search(search_key)
76 avl_search_end = time.perf_counter()
77 avl_search_time = avl_search_end - avl_search_start
78
79 # --- 4. Calculate Ratio ---
80 # Handle division by zero just in case
81 search_ratio = 0.0
82 if avl_search_time > 0:
83     search_ratio = bst_search_time / avl_search_time
84
85 # --- 5. Write Report ---
86 print(f"Writing analysis to {output_filename}...")
87 with open(output_filename, 'w') as report:
88     report.write(f"BST_HEIGHT: {bst_height}\n")
89     report.write(f"AVL_HEIGHT: {avl_height}\n")
90     report.write(f"BST_IS_BALANCED: {bst_balanced}\n")
91     report.write(f"AVL_IS_BALANCED: {avl_balanced}\n")
92     report.write(f"BST_SEARCH_TIME: {bst_search_time}\n")
93     report.write(f"AVL_SEARCH_TIME: {avl_search_time}\n")
94     report.write(f"SEARCH_RATIO_BST-vs-AVL: {search_ratio}\n")
95
96 print("Analysis complete.")

```

Submission

Submit your single Python script, `avl_vs_bst.py`, and the generated `analysis.txt` file.